Customer Retention for the Google Merchandise Store

Jay Patel | Sanchit Deora | Anshika Saxena | James Diffenderfer

Group 10

Objective

- Predicting ***Customer Retention*** for a company or store
- Given data collected on a Google Merchandise Store (GStore) customer, predict if that customer will return to shop at the GStore again

Software and Tools

Google Merchandise Store Data Set

- Data originally provided for **Kaggle Customer Revenue Prediction** competition
- Data provided in training and testing sets
  - **Training set (25 GB)**: User transactions from August 1, 2016 to April 30, 2018
  - **Testing set (8 GB)**: User transactions from May 1, 2018 to October 15, 2018
- List of **13 original features** (orange indicates JSON data)
  - fullVisitorId, channelGrouping, date, **device**, **geoNetwork**, **totals**, sessionId, socialEngagementType, **hits**, trafficSource, visitId, visitNumber, visitStartTime

Data Preprocessing Pipeline

- Created **customerReturns** labels
- **Subsampled rows** from full data set
  - Used **KMeans clustering** and **stratified sampling**
- Percentage of returning customers:
  - Full Dataset: 33.3 %
  - Subsampled Dataset: 33.4 %

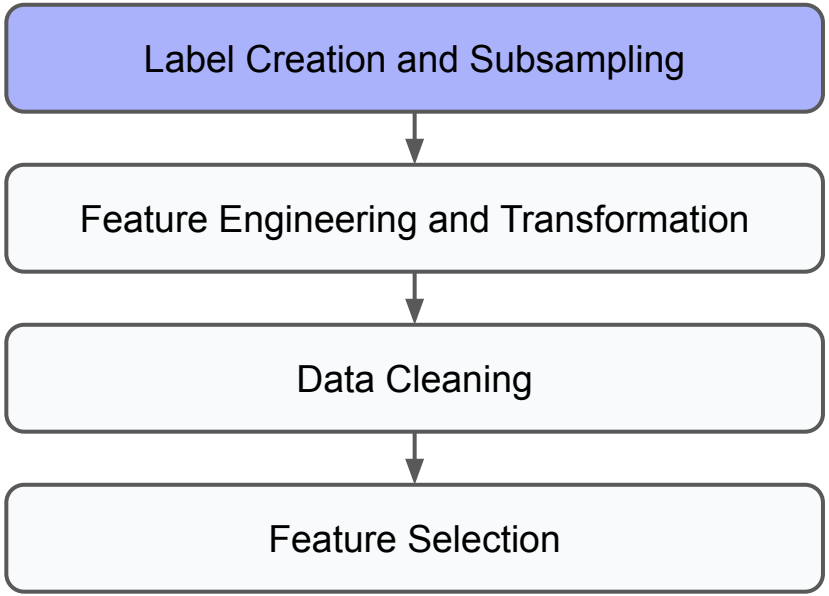

Data Preprocessing Pipeline

- Used data transformation to make data more comprehensive
  - Cyclic features (**Date**, **Time**)
  - Location features (**Latitude**, **Longitude**)
- Normalised data for better performance

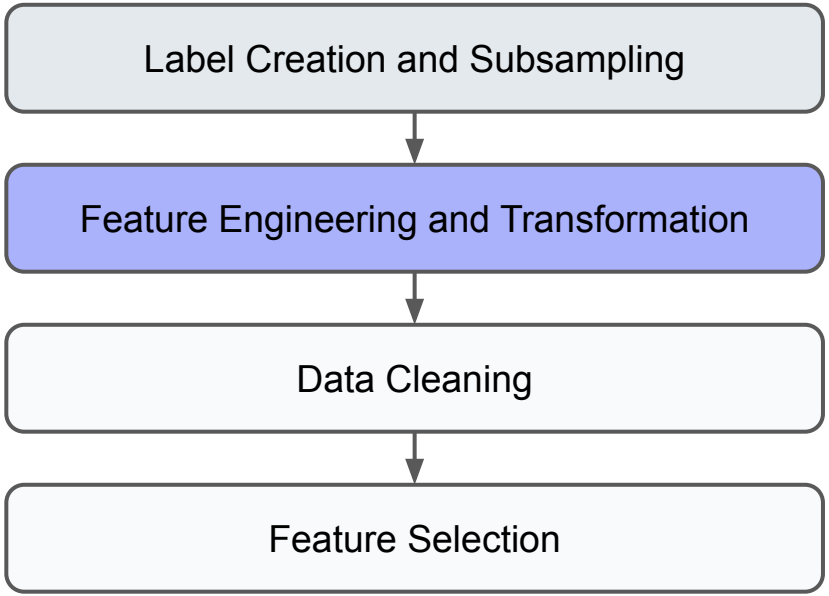

Data Preprocessing Pipeline

- Filled in **missing values**
  - Filled **-1** for numerical values
  - Filled **'UNK'** for string values
- **Deleted columns** if more than 90% of the data was missing

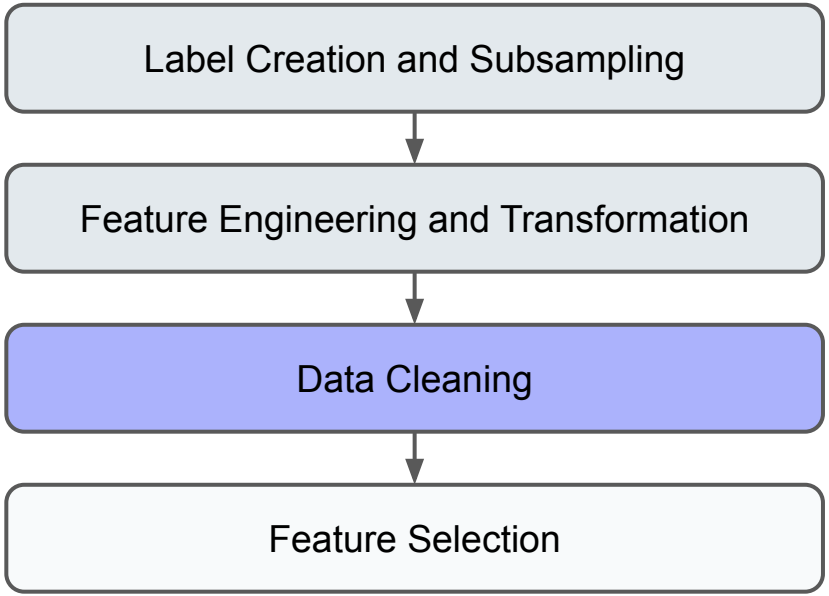

Data Preprocessing Pipeline

- Used **Extra Trees Classifier to get importance** for each feature
- Removed possibly non-contributing features

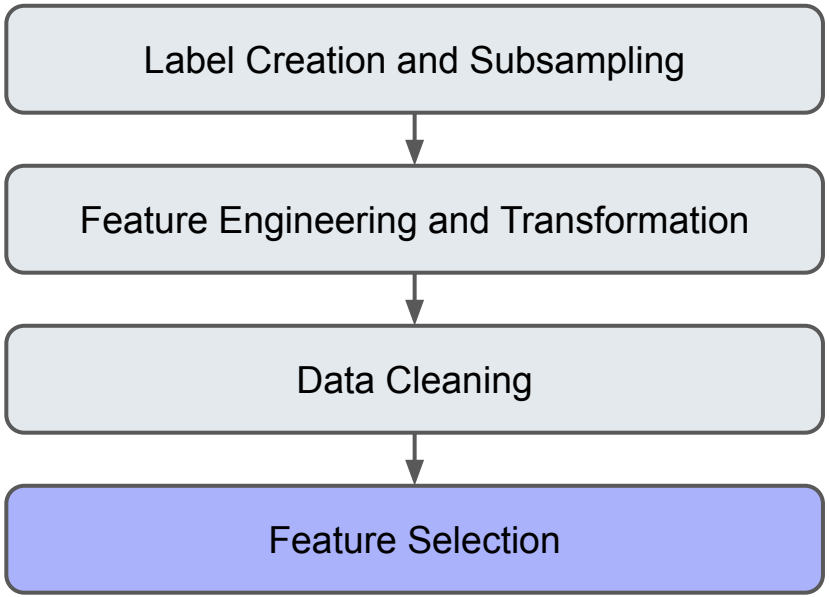

Data Set After Preprocessing

- Preprocessing pipeline resulted in **42 features**
- **Training/Testing split** is approximately **81/19**
  - **Training Set:** 1,537,503 samples
  - **Testing Set**: 361,429 samples

Modeling

- **Baseline Models**
  - Linear Regression
  - Gaussian Naive Bayes Classifier
  - Multinomial Naive Bayes Classifier
- **Trees**
  - Random Forest Classifier
  - XGBoosted Trees

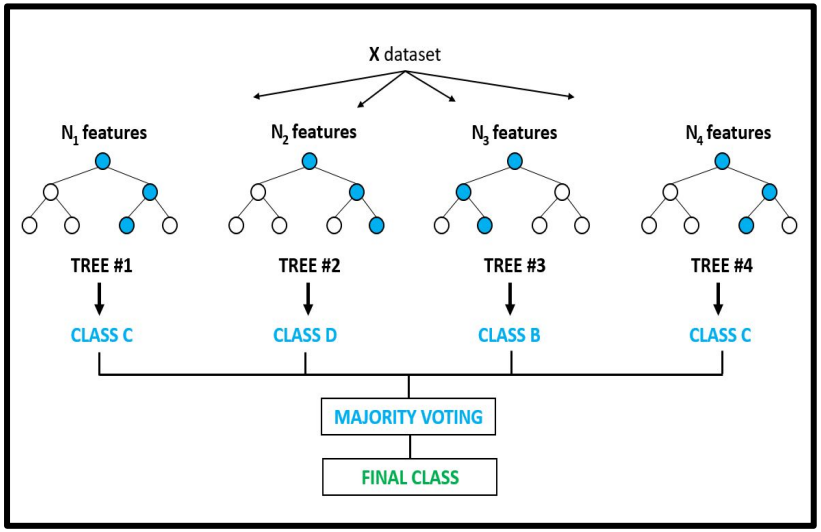

Random Forest

Modeling

- **Support Vector Machines**
  - Linear SVM
- **Neural Networks**
  - DNN: 2 hidden layers, ReLU activation, dropout layers
  - DNN: 1 hidden layer, dropout layers, batch normalization, ReLU activation
- **Clustering**
  - KMeans

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 48)	2016
batch_normalization (BatchNo	(None, 48)	192
activation (Activation)	(None, 48)	0
dropout (Dropout)	(None, 48)	0
dense_1 (Dense)	(None, 24)	1176
batch_normalization_1 (Batch	(None, 24)	96
activation_1 (Activation)	(None, 24)	0
dropout_1 (Dropout)	(None, 24)	0
dense_2 (Dense)	(None, 1)	25
batch_normalization_2 (Batch	(None, 1)	4
activation_2 (Activation)	(None, 1)	0

Total params: 3,509
Trainable params: 3,363
Non-trainable params: 146

DNN Model Summary

Model Accuracy

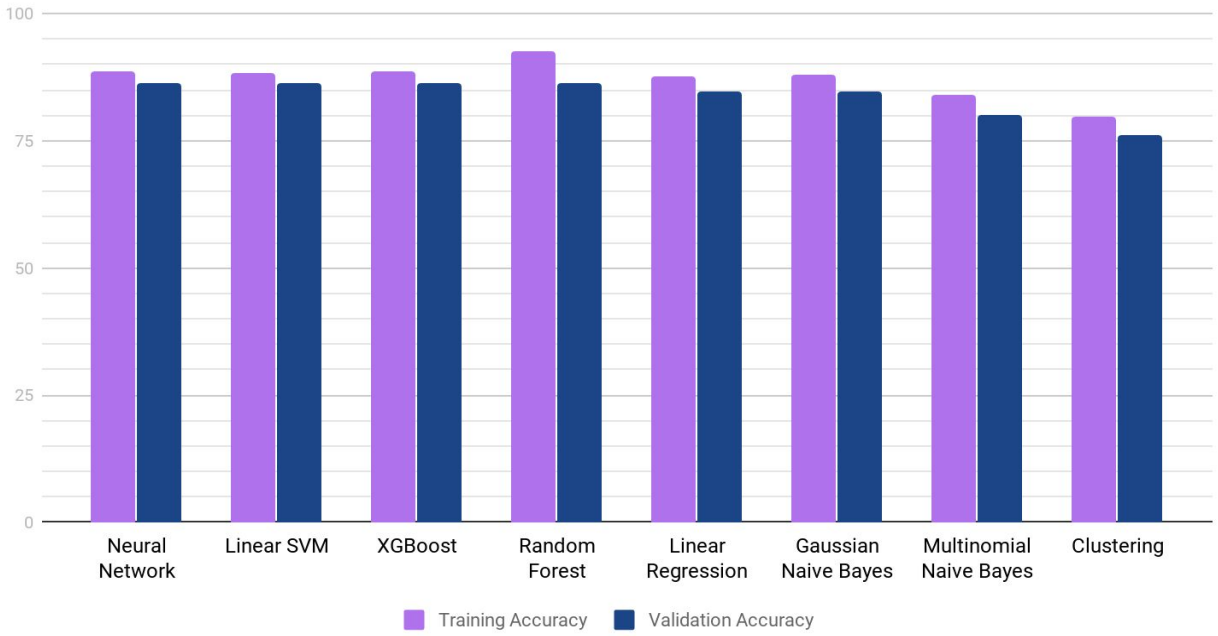

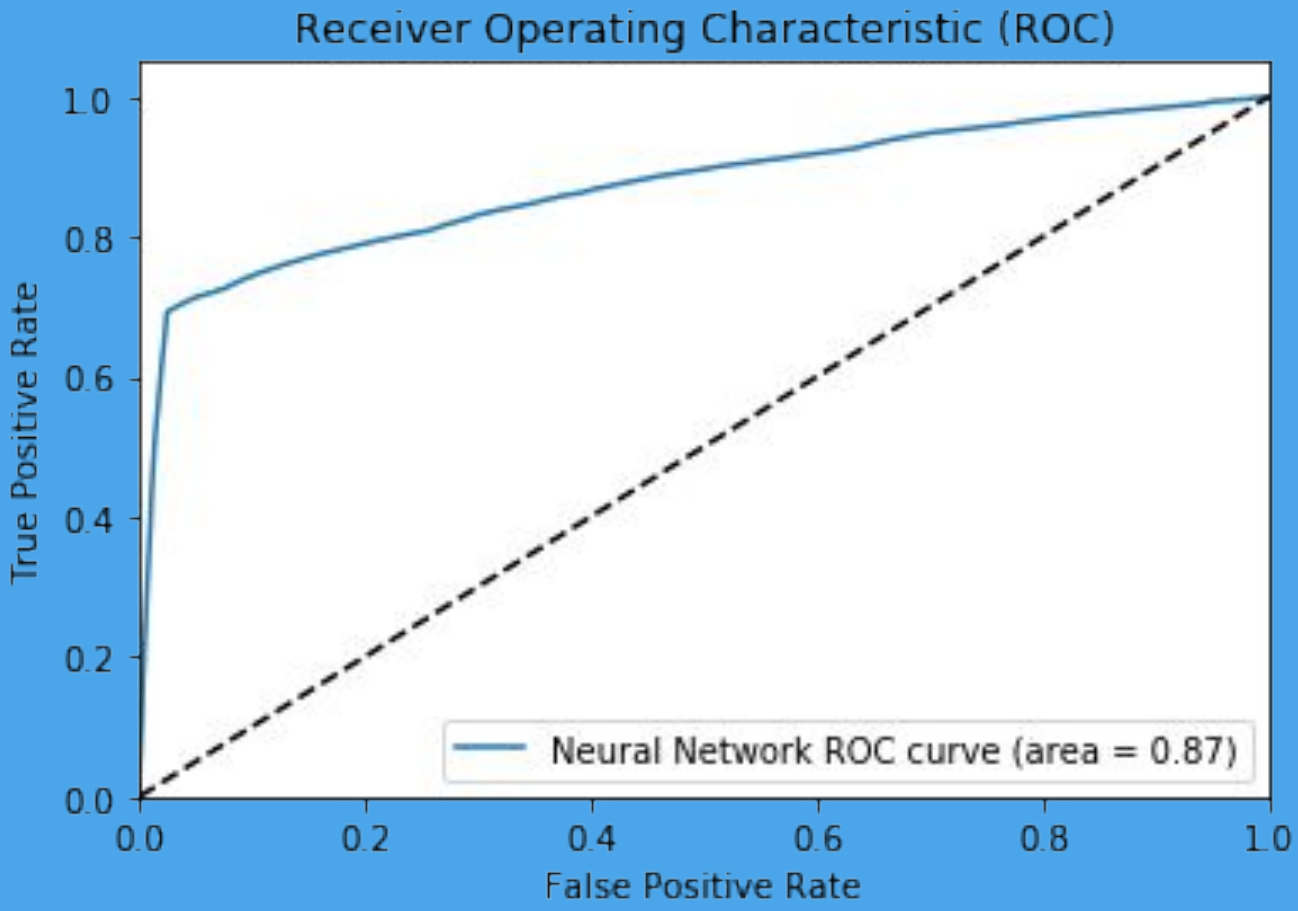
Receiver Operating Characteristic (ROC)
True Positive Rate
False Positive Rate
0.0
0.2
0.4
0.6
0.8
1.0
Neural Network ROC curve (area = 0.87)

Challenges

- **Data Preprocessing**
  - Large Data Set - Combined Training/Testing totals 33 GB
  - JSON columns in original data set
  - Processing missing values

- **Model Training**
  - Memory issues training certain models

Future Work

- **Ensemble Methods**
  - Combine models using weighted voting to create ensemble method
- **Develop Scalability**
  - Implement data preprocessing pipeline using Spark
- **Additional Data Preprocessing**
  - Attempt to extract and engineer more useful features from some JSON data

Thank you!